



Caché Process Memory

Version 2018.1
2024-11-07

Caché Process Memory

PDF generated on 2024-11-07

InterSystems Caché® Version 2018.1

Copyright © 2024 InterSystems Corporation

All rights reserved.

InterSystems®, HealthShare Care Community®, HealthShare Unified Care Record®, IntegratedML®, InterSystems Caché®, InterSystems Ensemble®, InterSystems HealthShare®, InterSystems IRIS®, and TrakCare are registered trademarks of InterSystems Corporation. HealthShare® CMS Solution Pack™, HealthShare® Health Connect Cloud™, InterSystems® Data Fabric Studio™, InterSystems IRIS for Health™, InterSystems Supply Chain Orchestrator™, and InterSystems TotalView™ For Asset Management are trademarks of InterSystems Corporation. TrakCare is a registered trademark in Australia and the European Union.

All other brand or product names used herein are trademarks or registered trademarks of their respective companies or organizations.

This document contains trade secret and confidential information which is the property of InterSystems Corporation, One Memorial Drive, Cambridge, MA 02142, or its affiliates, and is furnished for the sole purpose of the operation and maintenance of the products of InterSystems Corporation. No part of this publication is to be used for any other purpose, and this publication is not to be reproduced, copied, disclosed, transmitted, stored in a retrieval system or translated into any human or computer language, in any form, by any means, in whole or in part, without the express prior written consent of InterSystems Corporation.

The copying, use and disposition of this document and the software programs described herein is prohibited except to the limited extent set forth in the standard software license agreement(s) of InterSystems Corporation covering such programs and related documentation. InterSystems Corporation makes no representations and warranties concerning such software programs other than those set forth in such standard software license agreement(s). In addition, the liability of InterSystems Corporation for any losses or damages relating to or arising out of the use of such software programs is limited in the manner set forth in such standard software license agreement(s).

THE FOREGOING IS A GENERAL SUMMARY OF THE RESTRICTIONS AND LIMITATIONS IMPOSED BY INTERSYSTEMS CORPORATION ON THE USE OF, AND LIABILITY ARISING FROM, ITS COMPUTER SOFTWARE. FOR COMPLETE INFORMATION REFERENCE SHOULD BE MADE TO THE STANDARD SOFTWARE LICENSE AGREEMENT(S) OF INTERSYSTEMS CORPORATION, COPIES OF WHICH WILL BE MADE AVAILABLE UPON REQUEST.

InterSystems Corporation disclaims responsibility for errors which may appear in this document, and it reserves the right, in its sole discretion and without notice, to make substitutions and modifications in the products and practices described in this document.

For Support questions about any InterSystems products, contact:

InterSystems Worldwide Response Center (WRC)

Tel: +1-617-621-0700

Tel: +44 (0) 844 854 2917

Email: support@InterSystems.com

Table of Contents

Caché Process Memory	1
1 Introduction	1
2 Managing The Process Space	1
2.1 \$ZSTORAGE	1
2.2 Maximum Per-Process Memory	2
2.3 \$STORAGE	2
3 <STORE> Errors	2
3.1 Special Considerations	3

Caché Process Memory

1 Introduction

Processes use many different resources to accomplish their goals. These include some or all of CPU cycles, memory, external storage, network bandwidth, and others. This article is about memory usage. Specifically, it deals with the memory that Caché allocates for data storage, such as:

- Public And Private Variables

These are assigned memory space when a value is first assigned to them. In the case of local arrays, the combination of the local variable name plus the values of all the subscripts references a single variable value.

Variables containing strings shorter than 32,768 characters take up space associated with [\\$STORAGE](#). At the present time, long strings (those equal to or longer than 32KB) are stored differently and do not occupy space accounted for in [\\$STORAGE](#).

- Object Instances

Whenever an object is instantiated, space is allocated to hold the current contents of the object and possibly the objects it references. That space is returned when the last object reference is removed.

- Local I/O Buffers

Caché stores the I/O buffers associated with the devices in use by this process in the process space. For details, see the ObjectScript [OPEN](#) command and the [Caché I/O Device Guide](#).

2 Managing The Process Space

A process begins with an initial pool of memory to use for the entities described above. As the application creates them, they consume memory from the pool; when the application removes them, their memory is returned to the pool. For example, when a routine begins executing it almost always creates local variables that consume some memory; when the routine returns and those variables go out of scope, the memory used by those variables is returned and becomes available for re-use.

When an application requires memory, and the process does not have a sufficiently large (contiguous) area of memory available in its pool to satisfy the demand, the process requests an additional chunk of memory from the underlying operating system to add to its pool. Later, if that memory chunk becomes entirely unused, it will be returned to the operating system. Since the order in which entities are assigned memory and the order in which those entities are removed from memory are not necessarily mirror images of one another, as execution proceeds, memory becomes fragmented to some degree. This can affect the allocation and deallocation of memory from the operating system described above.

2.1 \$ZSTORAGE

A Caché process may use up to 2TB of memory. To aid in managing memory usage, Caché provides a way for an administrator or application to set a smaller limit on memory consumption. This value is stored in the system variable, [\\$ZSTORAGE](#) of each process so that [\\$ZSTORAGE](#) always contains the maximum allowed size of process memory (in KB).

The value of `$ZSTORAGE` is specified in units of 1KB. The minimum value allowed is 128, that is, 128KB. The maximum value a process can set for `$ZSTORAGE` is 2TB ($2^{31} * 1\text{KB}$) of memory. Attempts to set a value smaller than the minimum or larger than the maximum will default to the minimum or maximum, respectively.

2.2 Maximum Per-Process Memory

This value is set via the Management Portal. To access the relevant page, select **System Administration > Configuration > System Configuration > Memory and Startup**. On the page that comes up, set the value in `Maximum Per-Process Memory (KB)`.

In the configuration file (`cache.cpf`) this parameter is known as `bbsiz`. This value is the initial value for `$ZSTORAGE` when a process starts.

Note: The memory limit for a process can also be set when the process is started via the ObjectScript `JOB` command.

2.3 \$STORAGE

The system variable, `$STORAGE`, represents the amount of storage still available to the running process. It is given in bytes. When the process request for memory is larger than the value in `$STORAGE` or the request to allocate memory from the operating system fails, it generates a <STORE> error.

3 <STORE> Errors

When satisfying the process's request for memory would cause the value of `$STORAGE` to become negative, or the request to allocate memory from the operating system fails, it generates a <STORE> error. With regard to handling <STORE> errors where `$STORAGE` becomes negative, a Caché process can be thought of as being in one of two modes: normal mode and low-memory mode.

- Normal Mode

When a process is in normal mode and makes a request for memory that would otherwise cause `$STORAGE` to go negative, the process throws a <STORE> error and enters low-memory mode.

- Low-Memory Mode

In low-memory mode, operations are allowed to push `$STORAGE` negative, in order to allow some additional memory for the application to handle the error and clean up. When a process in low-memory mode frees memory such that the value of `$STORAGE` rises to at least 256KB (or 25% of `$ZSTORAGE` if that is lower) the process returns to normal mode. In low memory mode a lower limit of approximately -1MB is placed on `$STORAGE`. Any operation that would otherwise cause `$STORAGE` to fall below that limit results in a <STORE> error. The value of the lower limit is defined by the value of `$STORAGE` upon entering low-memory mode, minus 1MB.

Note: A process can set `$ZSTORAGE` to any value in its allowed range. If `$ZSTORAGE` is set to a value less than is currently in use, `$STORAGE` will have a negative value. If this occurs while the process is in normal mode, the next operation that allocates memory will cause the process to get a <STORE> error and enter low memory mode with the lower limit equal to that value minus 1MB. If this occurs while the process is already in low memory mode, the lower limit remains unaltered.

For <STORE> errors which result from exceeding the -1MB low-memory-mode limit, or that result from a failure to allocate memory from the operating system, the behavior of the process is unpredictable because there is so little memory available.

The process may be able to handle the error normally, the error handler may get a <STORE> error, or the error handler may not be able to be invoked and the process may halt.

An error handler may address the <STORE> error using one or more of these approaches:

- Abort the computation that caused the memory request, possibly freeing up any storage that the computation had obtained up until the <STORE> error occurred.
- Attempt to generate more available memory itself by getting rid of unneeded data.
- Perform whatever cleanup is necessary (such as closing open files), and terminate the program.
- Set the value of `$ZSTORAGE` to a larger value allowing the process to continue and request more memory in the future.

3.1 Special Considerations

3.1.1 Platforms

Most Caché instances run on systems that have less than 2TB available to allocate per Caché process. On such systems, when a Caché process exhausts the available system memory (real physical memory plus available swap space), the underlying system may deal with this condition in a number of ways. Some examples are:

- On some platforms, such as HPUX, the system will send a signal that causes the Caché process to terminate.
- On some platforms (for example Linux and AIX), the system uses a heuristic algorithm to kill the process that it deems to be the most offensive. This may be the Caché process, but it could also be another chosen process.
- Some systems deal with memory exhaustion by generating a kernel “panic” which crashes the underlying operating system.
- Some systems can handle the memory exhausted situation, but the recovery may result in access violations in the Caché process.

Good programming practice indicates that a Caché process should not depend on the error recovery algorithm used by the underlying platform. Instead, such processes should provide adequate design, testing, and logging to allow the process to estimate and manage its own resource requirements appropriately.

